

# Everything at Google runs in containers:

- Gmail, Web Search, Maps, ...
- MapReduce, batch, ...
- GFS, Colossus, ...
- Even GCE itself: VMs in containers

## We launch over **2 billion** containers **per week**.



Shipping Containers At Clyde, by Steve Gibson

# Kubernetes

Greek for *“Helmsman”*; also the root of the word *“Governor”*

- Container orchestration
- Runs Docker containers
- Supports multiple cloud and bare-metal environments
- Inspired and informed by Google’s experiences and internal systems
- **Open source**, written in **Go**

Manage applications, not machines



# Design principles

**Declarative > imperative:** State your desired results, let the system actuate

**Control loops:** Observe, rectify, repeat

**Simple > Complex:** Try to do as little as possible

**Modularity:** Components, interfaces, & plugins


**Legacy compatible:** Requiring apps to change is a non-starter

**No grouping:** Labels are the only groups


**Cattle > Pets:** Manage your workload in bulk

**Open > Closed:** Open Source, standards, REST, JSON, etc.

# High level design



# Primary concepts



**Container:** A sealed application package (Docker)

**Pod:** A small group of tightly coupled Containers

example: content syncer & web server

**Controller:** A loop that drives current state towards desired state

example: replication controller

**Service:** A set of running pods that work together

example: load-balanced backends

**Labels:** Identifying metadata attached to other objects

example: phase=canary vs. phase=prod

**Selector:** A query against labels, producing a set result

example: all pods where label phase == prod



# Pods



# Pods

**Small group** of containers & volumes

**Tightly** coupled

The atom of cluster scheduling & placement


Shared namespace

- **share IP** address & localhost

Ephemeral

- can die and be replaced

**Example: data puller & web server**



# Pod lifecycle

Once scheduled to a node, pods do not move

- restart policy means restart **in-place**

Pods can be observed *pending, running, succeeded, or failed*

- *failed* is **really** the end - no more restarts
- no complex state machine logic

Pods are **not rescheduled** by the scheduler or apiserver

- even if a node dies
- controllers are responsible for this
- keeps the scheduler **simple**

Apps should consider these rules

- Services hide this
- Makes pod-to-pod communication more formal



# Labels

Arbitrary metadata

Attached to any API object

Generally represent **identity**

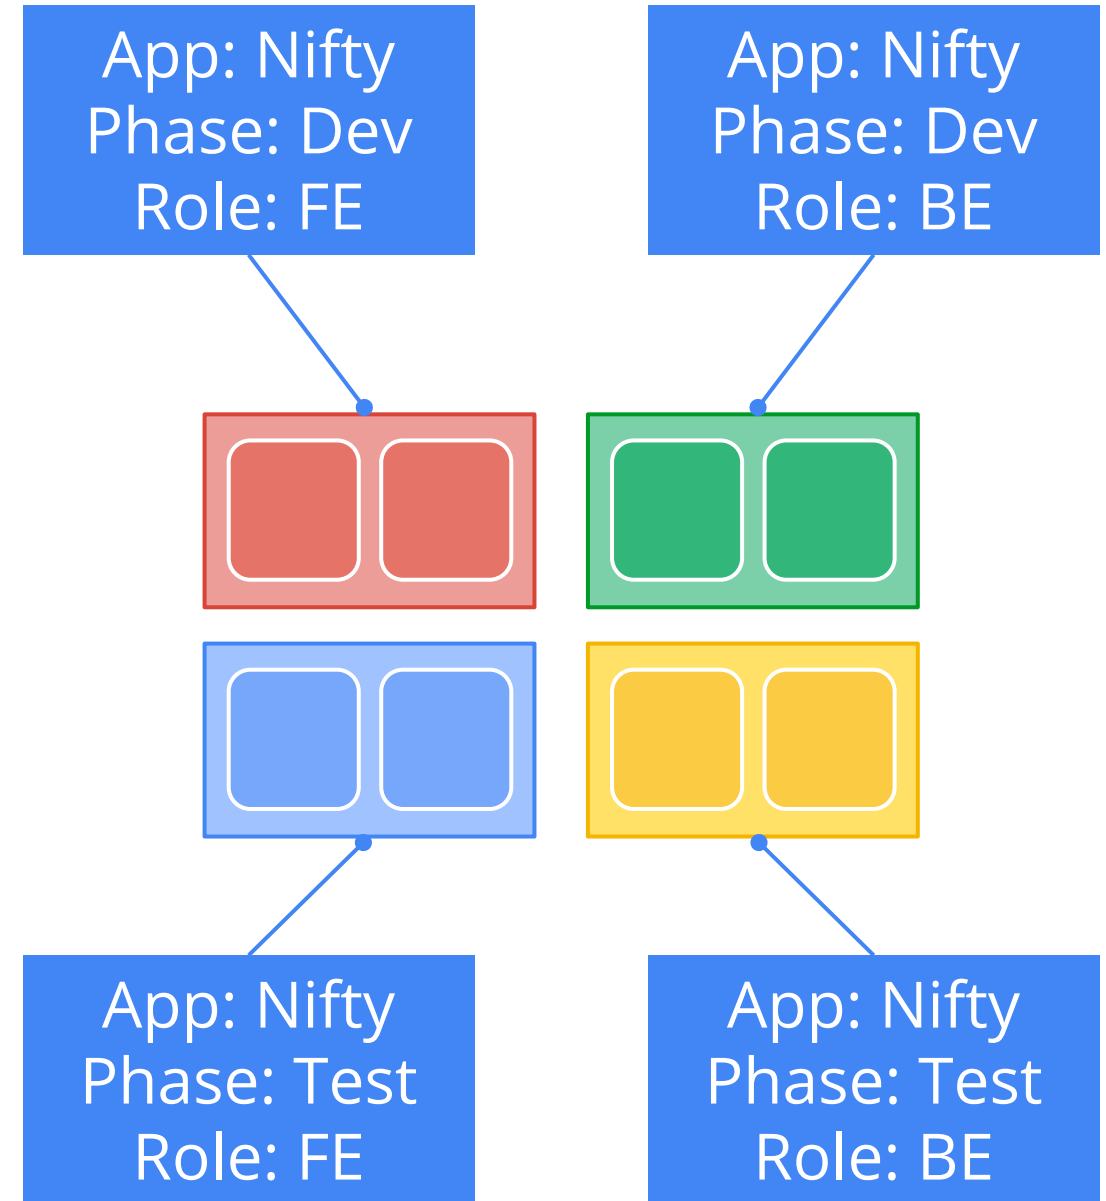
Queryable by **selectors**

- think SQL *'select ... where ...'*


The **only** grouping mechanism

- pods under a ReplicationController
- pods in a Service
- capabilities of a node (constraints)

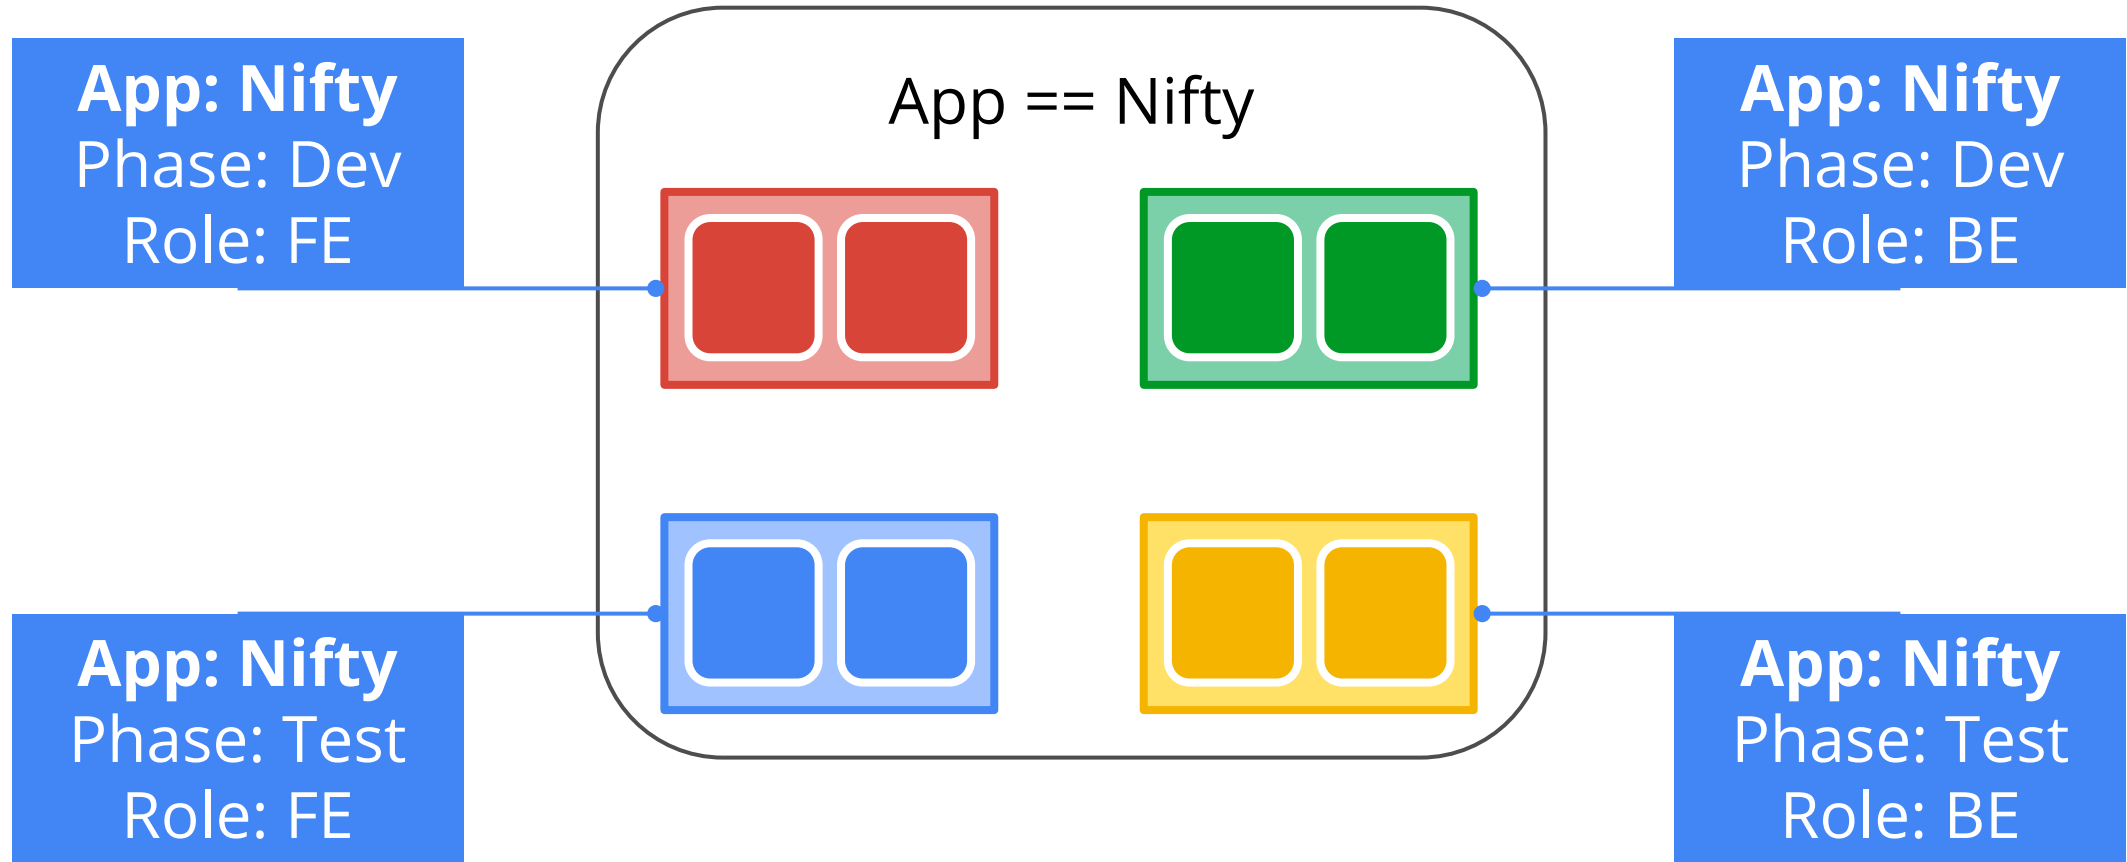
**Example: “phase: canary”**



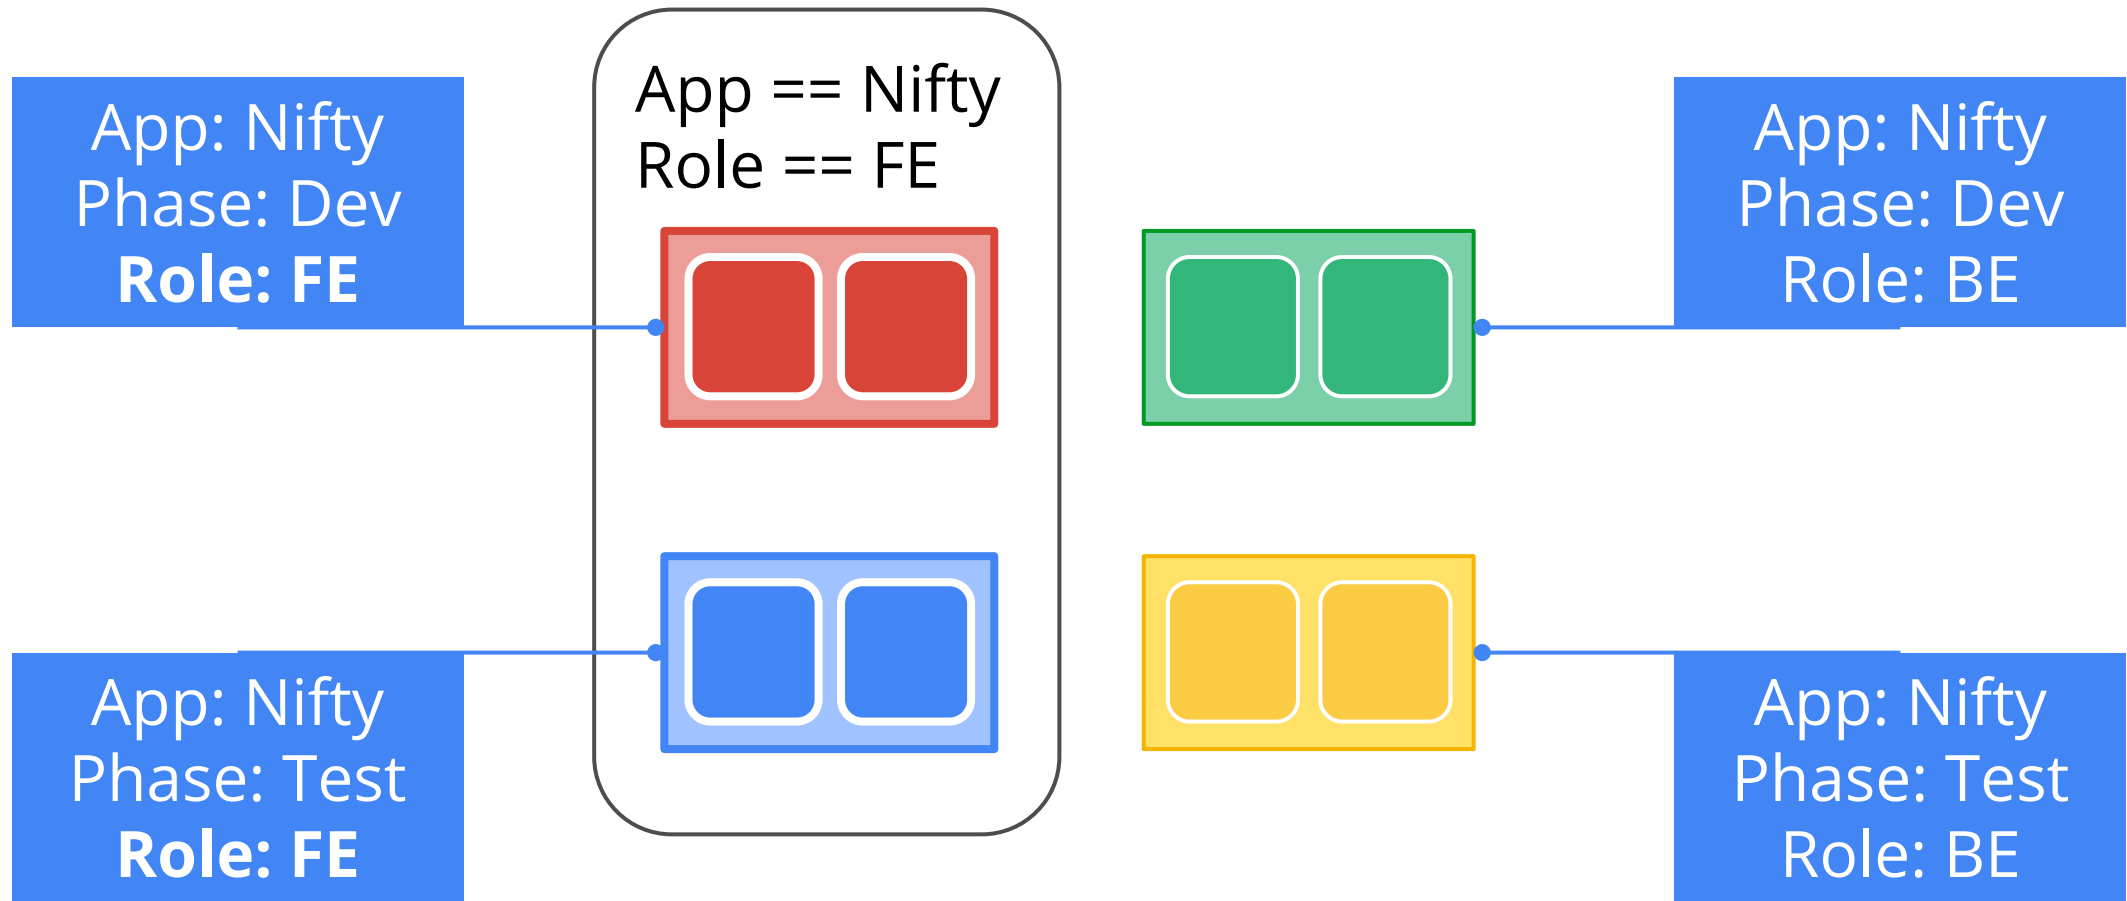
# Selectors



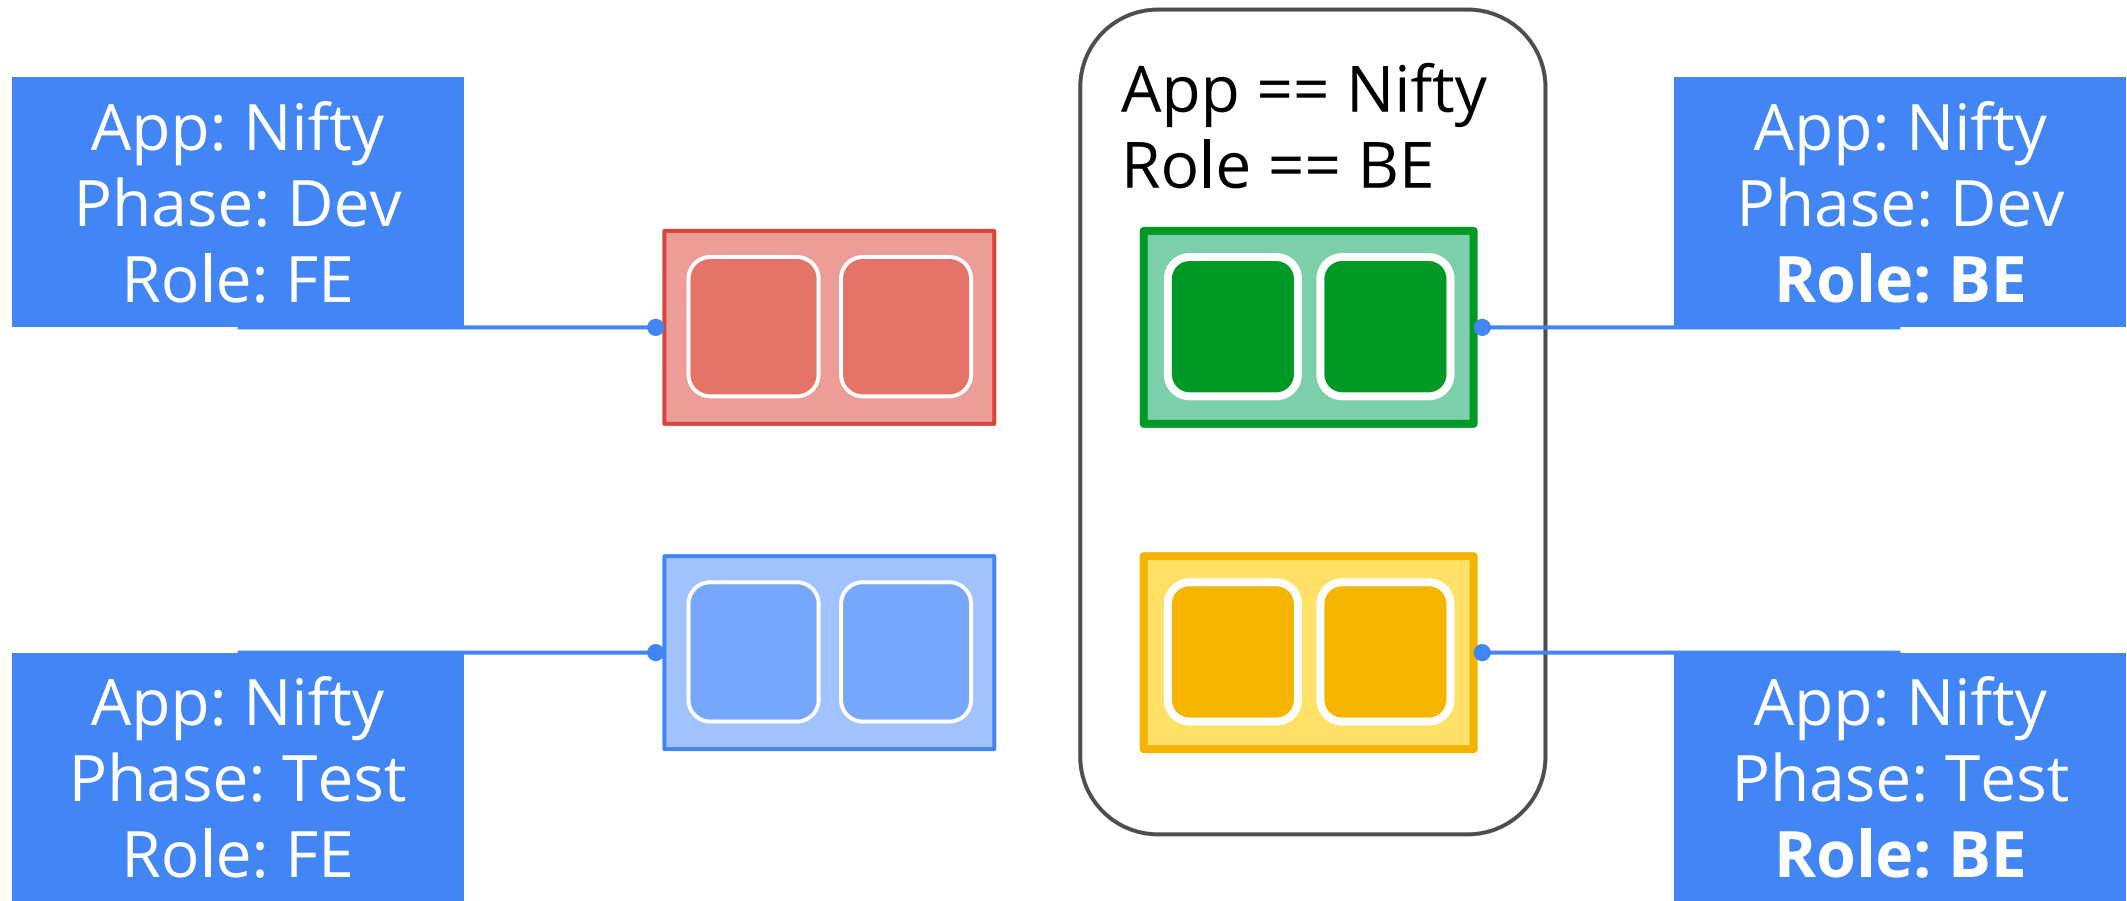
# Selectors



# Selectors

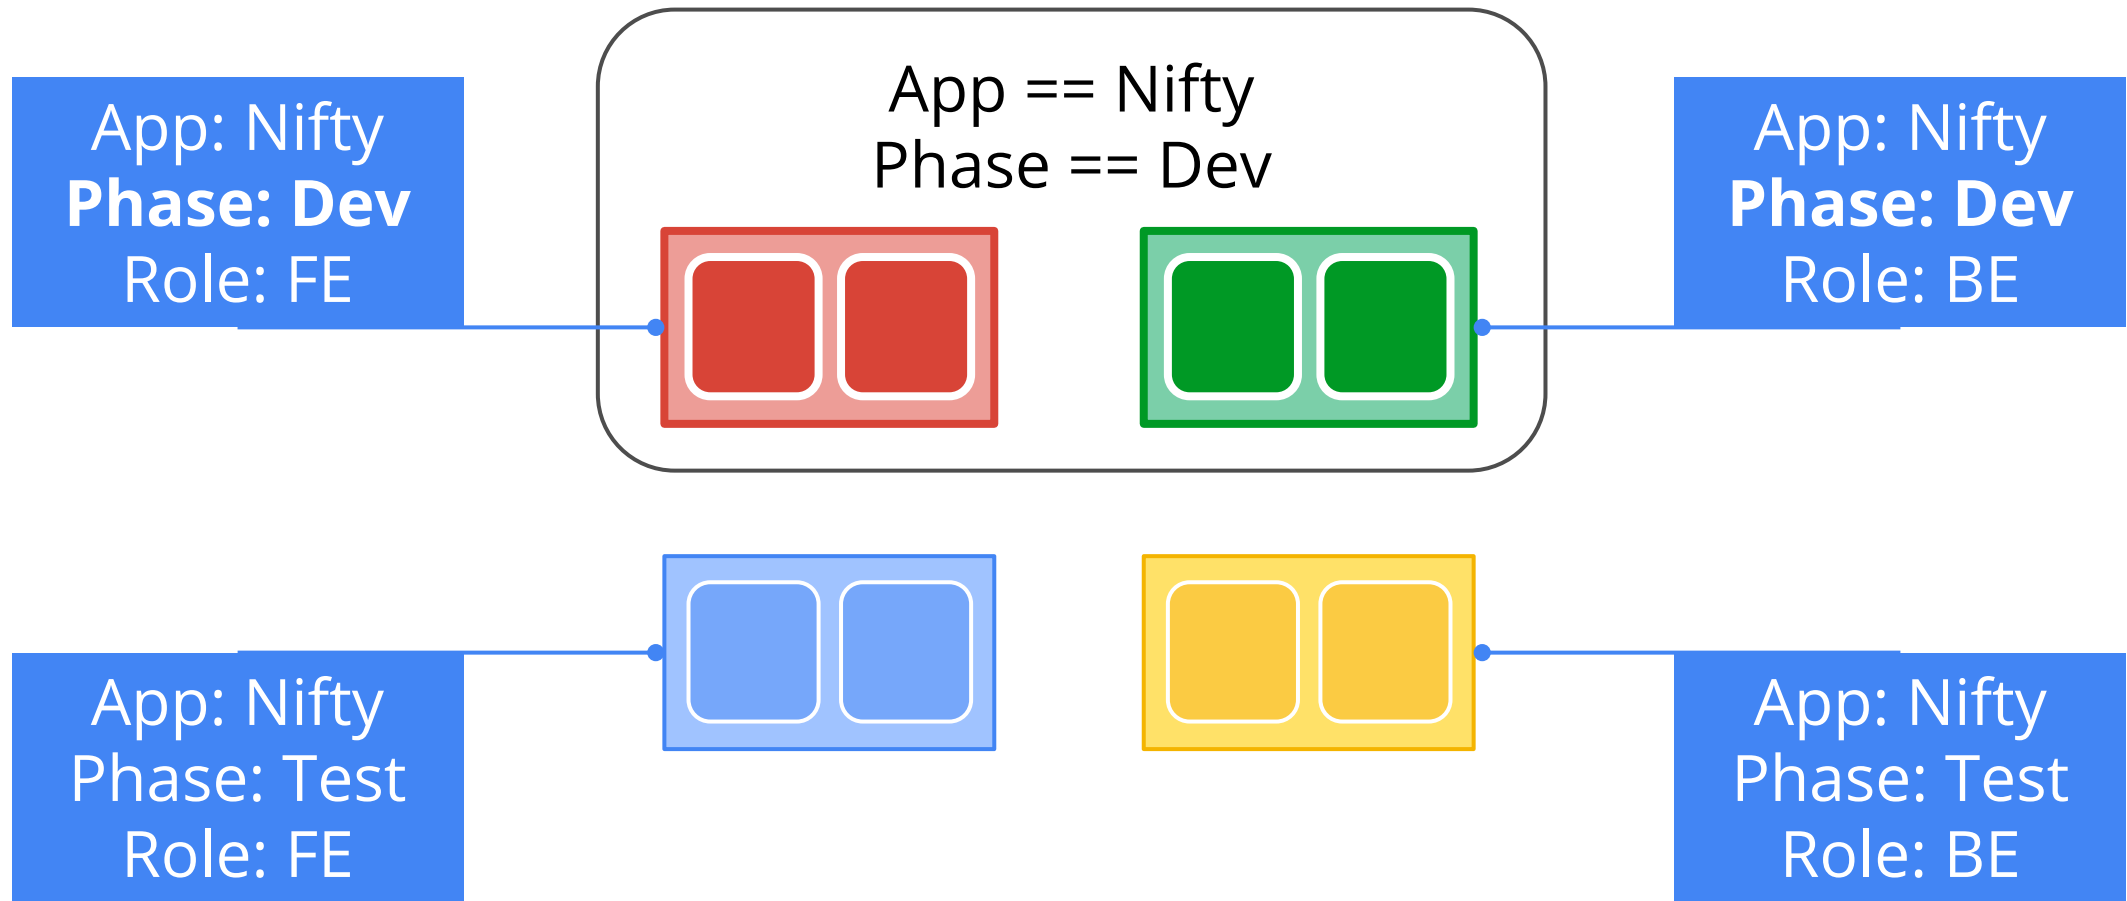


# Selectors

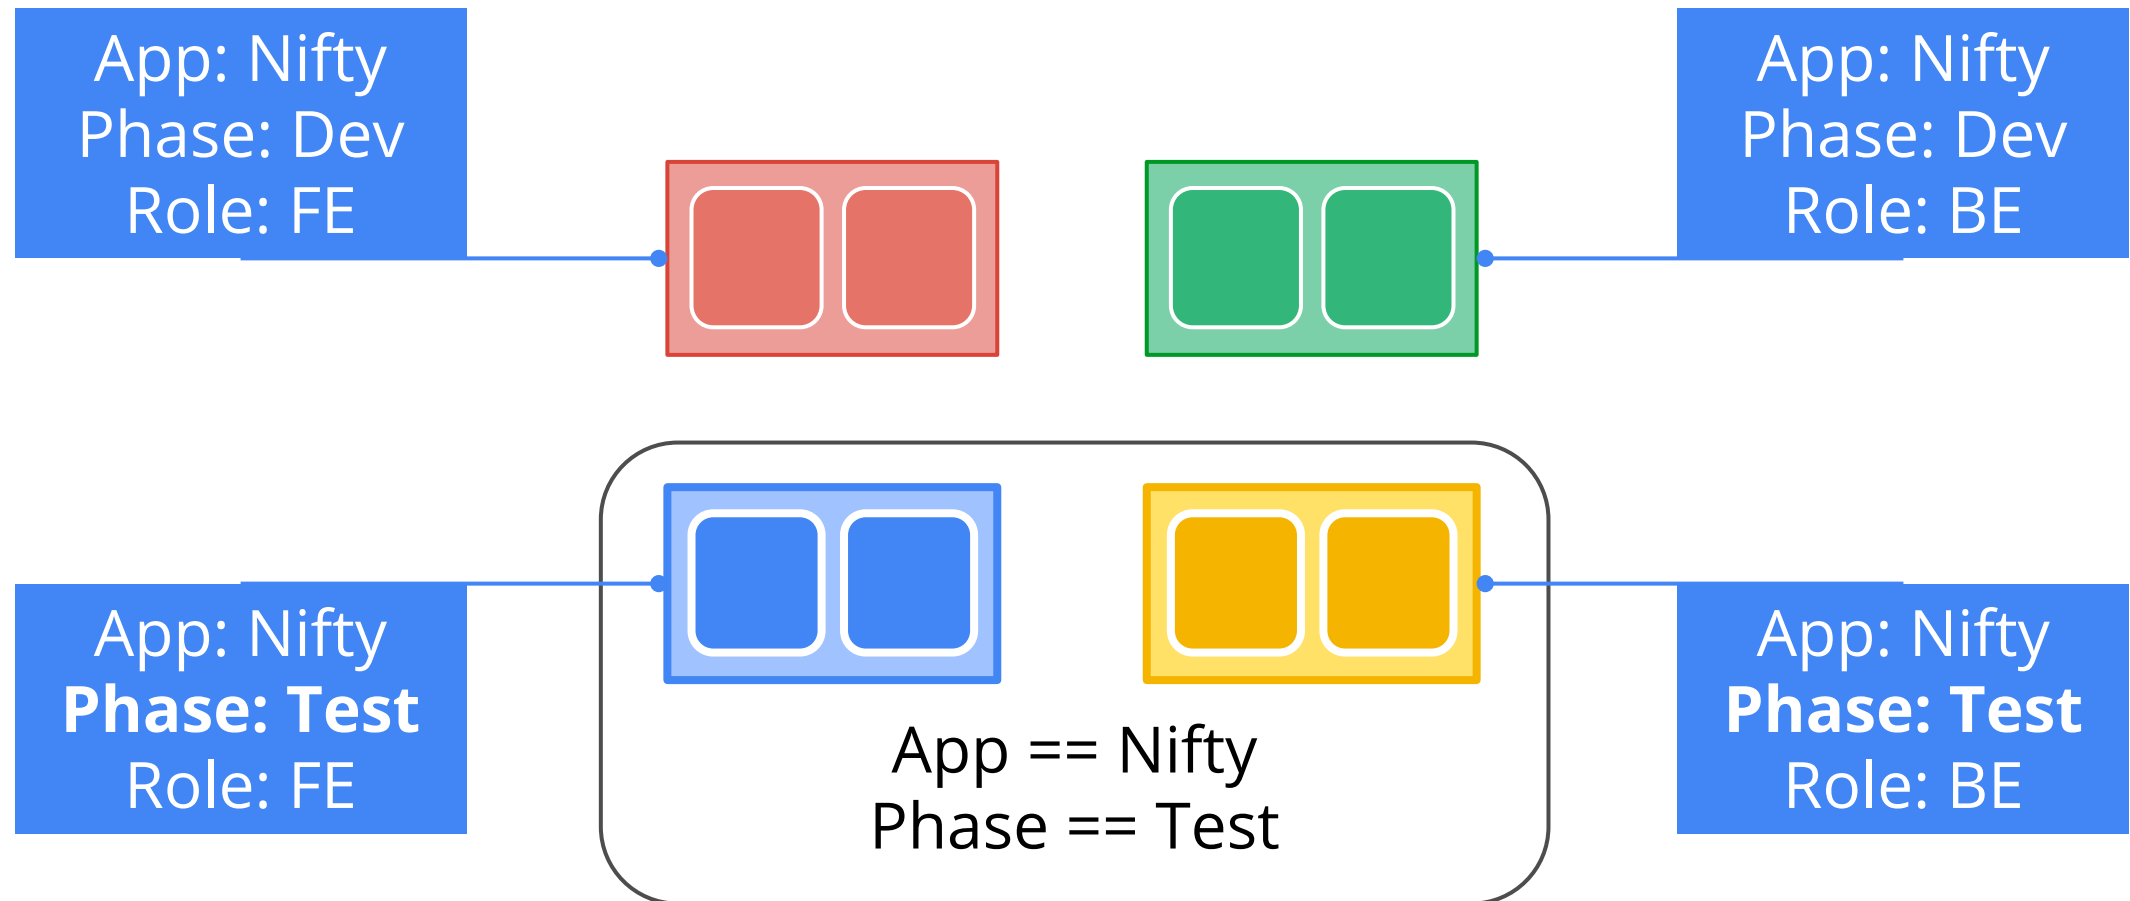




# Selectors



# Selectors



# Replication Controllers

Canonical example of control loops

Runs out-of-process wrt API server

Have 1 job: ensure N copies of a pod

- if too few, start new ones
- if too many, kill some
- group == selector

Cleanly layered on top of the core


- all access is by public APIs

Replicated pods are fungible

- No implied ordinality or identity

## Replication Controller





- Name = "nifty-rc"
- Selector = {"App": "Nifty"}
- PodTemplate = { ... }
- NumReplicas = 4




# Replication Controllers

## Replication Controller

- Desired = 4
- Current = 4






# Replication Controllers








# Replication Controllers



## Replication Controller

- Desired = 4
- **Current = 3**




# Replication Controllers



## Replication Controller

- Desired = 4
- Current = 4




# Replication Controllers



## Replication Controller

- Desired = 4
- **Current = 5**

# Replication Controllers



## Replication Controller

- Desired = 4
- Current = 4

# Pod networking

## Pod IPs are **routable**

- Docker default is private IP


## Pods can reach each other without NAT

- even across nodes

## No **brokering** of port numbers

## This is a **fundamental requirement**

- several SDN solutions





# Services

A group of pods that **act as one** == Service

- group == selector

Defines access policy

- only “load balanced” for now


Gets a **stable** virtual IP and port

- called the service *portal*
- also a DNS name

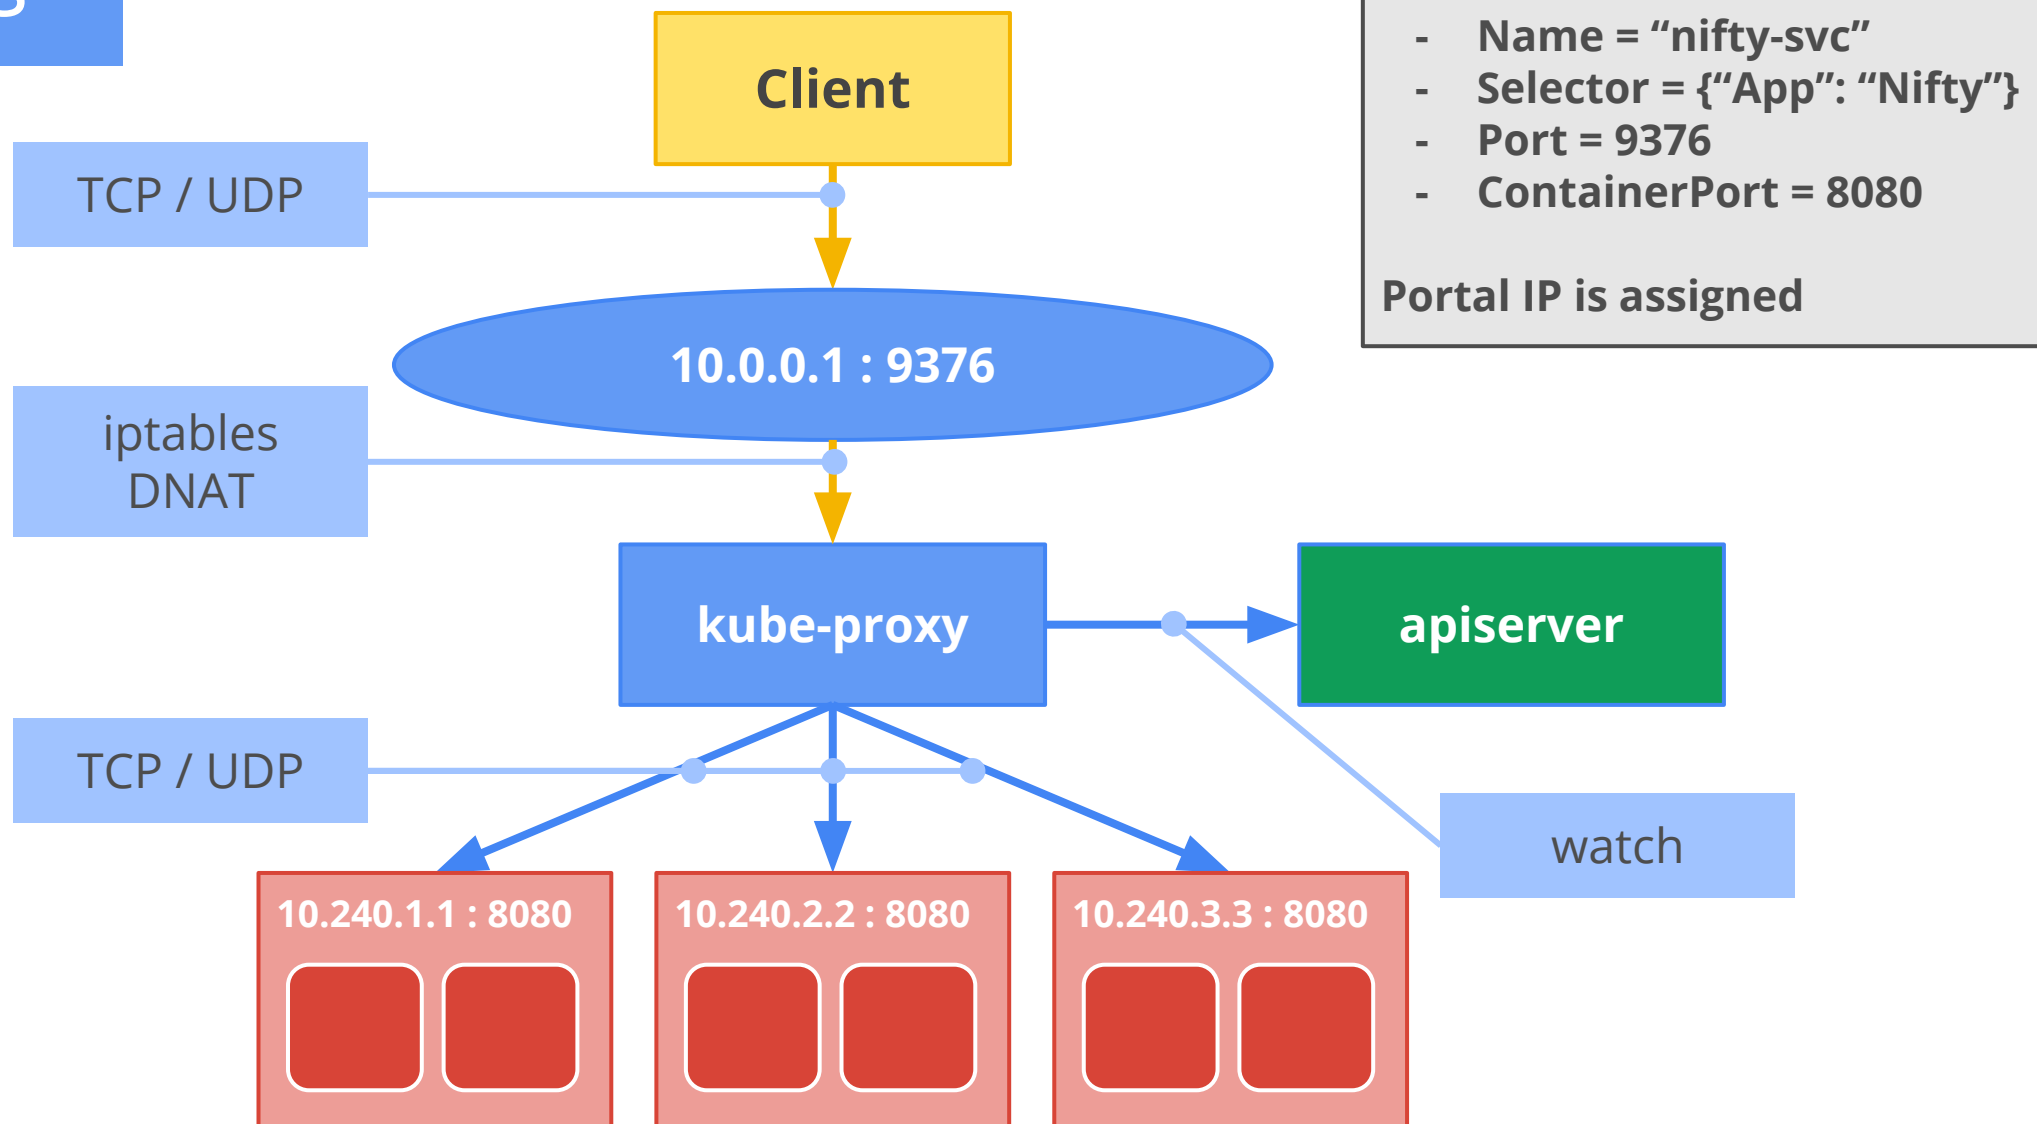
VIP is captured by *kube-proxy*

- watches the service **constituency**
- updates when backends change

Hide complexity - ideal for non-native apps



# Services



# Events

A central place for information about your cluster

- filed by any component: kubelet, scheduler, etc

Real-time information on the current state of your pod

- `kubectl describe pod foo`

Real-time information on the current state of your cluster

- `kubectl get --watch-only events`
- You can also ask only for events that mention some object you care about.

# Monitoring

Optional add-on to Kubernetes clusters

Run cAdvisor as a pod on each node

- gather stats from all containers
- export via REST

Run Heapster as a pod in the cluster

- just another pod, no special access
- aggregate stats

Run Influx and Grafana in the cluster

- more pods
- alternately: store in Google Cloud Monitoring



cAdvisor

# Logging

Optional add-on to Kubernetes clusters

Run fluentd as a pod on each node

- gather logs from all containers
- export to elasticsearch

Run Elasticsearch as a pod in the cluster

- just another pod, no special access
- aggregate logs

Run Kibana in the cluster

- yet another pod
- alternately: store in Google Cloud Logging



fluentd

